

Instructions for Authors of SBC Conferences Papers and Abstracts

Vitor Mendes paisante

¹Departamento de Ciência da Computação da Universidade Federal de Minas Gerais

{paisante}@dcc.ufmg.br

Abstract. *Alias analysis is one of the most fundamental techniques that compilers use to optimize languages with pointers. However, in spite of all the attention that this topic has received, the current state-of-the-art approaches inside compilers still face challenges regarding precision and speed. In particular, pointer arithmetic, a key feature in C and C++, is yet to be handled satisfactorily. This work presents a new alias analysis algorithm to solve this problem. The key insight of our approach is to combine alias analysis with symbolic range analysis. This combination lets us disambiguate fields within arrays and structs, effectively achieving more precision than traditional algorithms. To validate our technique, we have implemented it on top of the LLVM compiler. Tests on a vast suite of benchmarks show that we can disambiguate several kinds of C idioms that current state-of-the-art analyses cannot deal with. In particular, we can disambiguate 1.35x more queries than the alias analysis currently available in LLVM. Furthermore, our analysis is very fast: we can go over one million assembly instructions in 10 seconds..*

1. Context

Pointer analysis is one of the most fundamental compiler technologies. This analysis lets the compiler distinguish one memory location from others; hence, it provides the necessary information to transform code that manipulates memory. Given this importance, it comes as no surprise that pointer analysis has been one of the most researched topics within the field of compiler construction[?]. This research has contributed to make the present algorithms more precise [?, ?], and faster [?, ?]. Nevertheless, one particular feature of imperative programming languages remains to be handled satisfactorily by the current state-of-the-art approaches: the disambiguation of pointer intervals.

2. Problem

Mainstream compilers still struggle to distinguish intervals within the same array. In other words, state-of-the-art pointer analyses often fail to disambiguate regions addressed from a common base pointer via different offsets, as explained by Yong and Horwitz [?]. Figure 1 shows an example that state-of-the-art pointer analyses tend not to deal with satisfactorily. In the example, both stores on the *r* array possess different offsets and they do not alias. To figure out that such offsets are disjoint it is necessary to verify their ranges and add a new layer of analysis to the current pointer analyses available.

Field-sensitive pointer analysis, provide a partial solution to this problem. These analyses can distinguish different fields within a record, such as a struct in C [?], or a

```

1:  char* duplicate (int size , char* v) {
2:      if (size > 0) {
3:          char* r = malloc (size*2);
4:          int i;
5:          for (i = 0; i < size; i++) {
6:              r[i] = v[i];
7:              r[i+size] = v[i];
8:          }
9:      }
10:     else return NULL;
11: }

```

Figure 1. Example that state-of-the-art pointer analyses handle unsatisfactorily.

class in Java [?]. However, they rely on syntax that is usually absent in the low level program representations adopted by compilers. Shape analyses [?, ?] can disambiguate subparts of data-structures such as arrays, yet their scalability remains an issue to be solved. Consequently, many compiler optimizations, such as loop transformations, tiling, fission, skewing and interchanging [?, Ch.09], are very limited in practice. Therefore, we claim that, to reach their full potential, compilers need to be provided with more effective alias analyses.

3. Solution

This work describes such an analysis. We introduce an abstract domain that associates pointers with symbolic ranges. In other words, for each pointer p we conservatively estimate the range of memory slots that can be addressed as an offset of p . We let $\triangleright(p)$ be the global abstract address set associated with pointer p , such that if $loci + [l, u] \in \triangleright(p)$, then p may dereference any address from $@(loci) + l$ to $@(loci) + u$, where $loci$ is a program site that contains a memory allocation call, and $@(loci)$ is the actual return address of the *malloc* at runtime. We let $\{l, u\}$ be two *symbols* defined within the program code. Like the vast majority of pointer analyses available in the compiler literature, from Andersen’s work [?] to the more recent technique of Zhang *et al.* [?], our method is correct if the underlying program is also correct. In other words, our results are sound with respect to the semantics of the program if this program has no undefined behavior, such as out-of-bounds accesses.

The key insight of our research is the combination of pointer analysis with range analysis on the symbolic interval lattice. In a symbolic range analysis, ranges are defined as expressions of the program’s symbols, a symbol being either a constant or the name of a variable. There exist many approaches to symbolic range analyses in the literature [?, ?, ?]. The algorithms that we present in this work do not depend on any particular implementation. Nevertheless, the more precise the range analysis that we use, the more precise the analysis facts that we produce. In this work we have adopted the symbolic range analysis proposed in 1994 by William Blume and Rudolf Eigenmann [?].

4. Summary of experimental results

To validate our ideas, we have implemented them in the LLVM compilation infrastructure [?]. We have tested our pointer analysis onto three different benchmarks used in previous work related to pointer disambiguation: Prolangs [?], PtrDist [?] and Malloc-Bench [?]. As we show in Chapter ??, our analysis is linear on the size of programs. It can go over one-million assembly instructions in approximately 10 seconds. Furthermore, we can disambiguate 1.35x more queries than the alias analysis currently available in LLVM.

5. Summary of publications

The analysis described here was published on the International Symposium on Code Generation and Optimization (CGO) of 2016 held in Barcelona, Spain [?].

The technology behind it is also present on two other publications [?, ?]. In these papers, the algorithm presented here was used to infer the layout and the content of buffers transferred through the network. This was useful for verifying, in a safe communication line, if the information transferred between two programs through a network should be considered potentially dangerous or not. If proven not dangerous, guards for checking integer overflows may not be necessary. These articles proposed methods for such verification to be used on the internet of things (IoT), where simple devices could run significantly faster with a reduced number of integer overflow checks. The layout and content inference analysis used in these papers differed from our current approach by using a numerical range analysis, since a symbolic approach would not be of relevance for such application.

6. Limitations

Using the symbolic range analysis for comparing offsets is inherently imprecise. Issues lie on comparing two symbolic expressions and its difficulties, especially on comparing lower bounds with upper bounds from two different variables, and on local relationships of variables with overlapping ranges.

Figure 2 shows the first issue. Analyzing this algorithm, the symbolic range analysis returns the following ranges: $R(\sigma_a) = [a, b - 1]$, $R(\sigma_b) = [a + 1, b]$. This hinders the disambiguation of the two array accesses, $V[\sigma_a]$ and $V[\sigma_b]$, since it is impossible to prove that the ranges of σ_a and σ_b do not overlap. To do so, it would be necessary for the valuation of symbolic expressions to be able to say that $b - 1 < a + 1$ or that $b < a$, which it cannot do. So in this example, even though it is obvious that the two array accesses cannot alias to the same location since their offsets are obligatorily different by the **if** condition, we cannot disambiguate them.

Figure 3 shows the second issue. It again shows two array accesses that are obviously disjoint that our analysis cannot disambiguate. Analyzing this algorithm, the symbolic range analysis returns the following ranges: $R(i) = [0, 9]$, $R(j) = [1, 10]$. It's clear that these ranges overlap and that our analysis could not say that they are disjoint even though, at the array accesses, j is always different and greater than i .

These examples show two very interesting limitations of our analysis. They both can disable optimizations such as automatic parallelization of code and loop invariant code motion on some cases in which these optimizations might be desirable. It is clear that

```

1: ...
2: if  $a < b$  then
3:    $\sigma_a = \sigma(a)$ 
4:    $\sigma_b = \sigma(b)$ 
5:    $V[\sigma_a] \leftarrow \bullet$ 
6:    $V[\sigma_b] \leftarrow \bullet$ 
7:   ...
8: end if
9: ...

```

Figure 2. Example of our first described limitation. Even though we cannot prove that the symbolic ranges of σ_a and σ_b do not overlap, it is obvious that σ_a is always different from σ_b because of the if condition.

```

1: ...
2:  $i = 0$ 
3: while  $i < 10$  do
4:    $j = i + 1$ 
5:    $V[i] \leftarrow \bullet$ 
6:    $V[j] \leftarrow \bullet$ 
7:   ...
8:    $i++$ 
9: end while
10: ...

```

Figure 3. Example of our second described limitation. Even though we cannot prove that the ranges of i and j do not overlap, it is obvious that j is always different from i because it is always greater.

simply verifying if two offset ranges are disjoint is not enough to achieve ideal precision, since they can hide relationships between variables that are always true when the two variables are alive and that can help in disambiguating two pointers, even if there is a relationship between these pointers or between the offsets that constitute such pointers.

7. Future Work

The reason for the limitations exposed before lays on the fact that our analysis use range intervals to disambiguate pointers. In the examples of figure 2 and figure 3, the ranges of integer variables might either overlap or aren't comparable. But, in the examples, variables do have relationships between them that allow for disambiguation in the form of inequalities ($a < b$ and $i < j$). To take advantage of such relationships we intend to use a lattice similar to Pentagons.

Pentagons is an abstract domain invented by Logozzo and Fähndrich to infer symbolic bounds to the integer variables used in programs [?, ?]. This abstract domain is formed by the combination of two lattices. The first lattice is the *integer interval domain* [?], which maps integer variables to ranges $[l, u]$ of numeric lower (l) and upper (u) bounds. The second lattice is the *strict upper bound*, which maps each variable v to a set $L_<$ of other variables, so that if $u \in L_<(v)$, at a given program point p , then $u < v$ at p .

Since their debut [?], Pentagons have been used in several different ways. For

instance, Logozzo and Fähndrich have employed this domain to eliminate array bound checks in strongly typed programming languages [?], and to ensure absence of division by zero or integer overflows in programs. Moreover, Nazaré *et al.* [?] have used Pentagons to reduce the overhead imposed by AddressSanitizer [?] to guard C against out-of-bounds memory accesses. The appeal of pentagons comes from two facts. First, this abstract domain can be computed efficiently – in quadratic time on the number of program variables. Second, as an enabler of compiler optimizations, Pentagons have been proven to be substantially more effective than other forms of abstract interpretation of similar runtime [?].

Our initial and recent use of pentagons for such purpose has been successful. It is able to handle programs as large as SPEC’s gcc in a few minutes and go through SPEC CPU 2006 [?] in able time. It has proven very useful in some cases, such as SPEC 470.lbm. In future work, we plan to investigate better splitting strategies and other more expressive lattices to improve the global precision of our analyses.

8. Final Conclusions

In this work we have presented a new alias analysis technique that handles, within the same theoretical framework, the subtleties of pointer arithmetic and memory indexation. Our technique can disambiguate regions within arrays and C-like structs using the same abstract interpreter. We have achieved precision in our algorithm by combining alias analysis with classic range analysis on the symbolic domain. Our analysis is fast, and handles cases that the implementations of pointer analyses currently available in LLVM cannot deal with.

Apart from this contribution, there is plenty to study on the area of pointer analysis, and the area of pointer arithmetics still needs quite a bit of research. Their dire needs are very efficient static analyses that run fast on very big programs, and very lean dynamic analyses. Our focus has been on the static analysis side. Lazy implementations, where main computations are made on the query moment, seem to be a very promising take on alias analyses and can expedite runtime. Focus on integrating new proposals to bigger compilation frameworks and existing optimizations should also be explored by the community on an effort of making new technology more usable across researchers and projects. There is still a lot of work to be done and we hope that our contribution can be only a building block of a much bigger effort from many more scientists.

Referências